
NCURA
Regional Meeting '26
Terms and Conditions May Apply
Using AI to Translate Award Agreements into PI-Ready Action
Fable Mangold | CSU Channel Islands

REVEAL THE STRUCTURE



KEEP THE EXPERTISE.

Reimagine the workflow.



Terms and Conditions *May Apply*

Using AI to Translate Award Agreements into PI-Ready Action

FABLE MANGOLD / CSU CHANNEL ISLANDS



Use the AI environment your institution approves.

START WITH POLICY

Use institutionally approved accounts and tools.
Follow ITS guidance on AI, privacy, security, and permitted data.
Keep human review and approval in the workflow.

ADAPT THE ARCHITECTURE

ChatGPT: Projects or custom GPTs
Gemini: Gems
Claude: Projects
Agent seeds need platform-specific adaptation and testing.

The method is portable in principle. The data permissions and implementation are not automatic.

An award arrives.
The work begins.

01 Review the terms

02 Identify obligations

03 Prepare the handoff

A familiar process. A surprising opportunity.

AI could read the agreement.

But could I trust the answers?

The document

What does the agreement actually say?

The response

Plausible language is not always supported.

The reviewer

Can I verify the claim efficiently?

“I don’t know.”

An acceptable answer. A better starting point.

Where did that answer come from?

> *“Some details were provided without citation. Can you explain why this is the case…?”*
> *Historical BeachSMART conversation, August 12, 2025.*

≡

[Doc]

Direct evidence

≈

[Interp]

Interpretation

—

[Absent]

Not located

⌂

[Internal]

Campus records

Source labels prompt verification. They do not prove correctness.

One question became a process.

01

SCOPE

What matters?

02

EVIDENCE

Where is it?

03

UNCERTAINTY

What is missing?

04

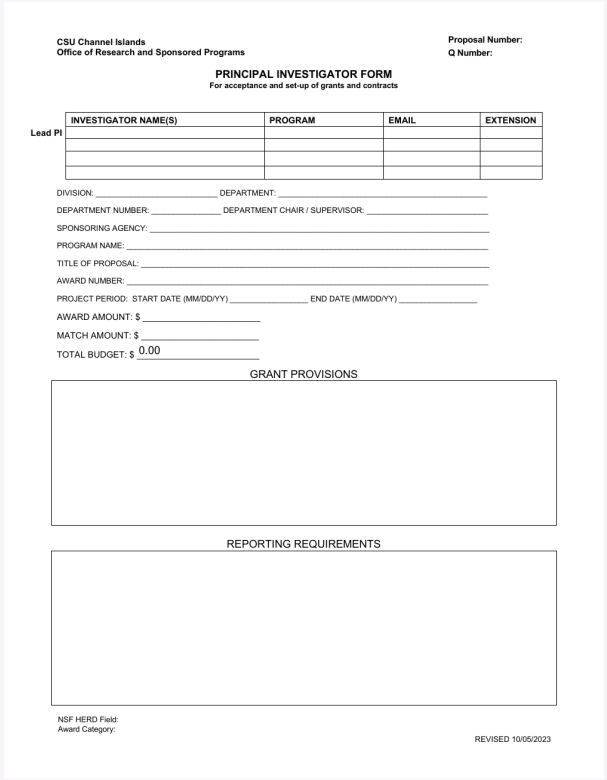
HUMAN

What matters next?

Next: translating the existing PI checklist into reusable instructions.

I didn’t start with AI.

A campus award-setup workflow already existed.



CSU Channel Islands
Office of Research and Sponsored Programs

Proposal Number:
Q Number:

PRINCIPAL INVESTIGATOR FORM
For acceptance and set-up of grants and contracts

	INVESTIGATOR NAME(S)	PROGRAM	EMAIL	EXTENSION
Lead PI				

DIVISION: _____ DEPARTMENT: _____
DEPARTMENT NUMBER: _____ DEPARTMENT CHAIR / SUPERVISOR: _____
SPONSORING AGENCY: _____
PROGRAM NAME: _____
TITLE OF PROPOSAL: _____
AWARD NUMBER: _____
PROJECT PERIOD: START DATE (MM/DD/YY) _____ END DATE (MM/DD/YY) _____
AWARD AMOUNT: $ _____
MATCH AMOUNT: $ _____
TOTAL BUDGET: $ 0.00

GRANT PROVISIONS

REPORTING REQUIREMENTS

NSF HERD Field:
Award Category:

REVISED 10/05/2023

01 Award agreement

Read the governing documents.

02 Administrative review

Interpret terms and identify missing information.

03 PI Checklist + meeting

Communicate selected obligations and hand off.

What if I could explain how I do this work?

A professional procedure became a reusable set of instructions.

“Nova, run the PI extraction.”

FOLLOW THE FORM

Use the PI Checklist order.

SHOW YOUR SOURCES

Classify and cite extracted information.

FLAG THE GAPS

Allow missing or internal-only answers.

WRITE FOR PEOPLE

Use plain language and separate responsibilities.

Where did that answer come from?

February 2026: a useful answer still needed more precise evidence.

FIRST OUTPUT

Technical reports due quarterly and a final report due after the project. [Doc]

Linked file, but no specific section or page in the displayed citation.

AFTER HUMAN FEEDBACK

Technical reports: quarterly, within 30 days; final within 60 days.

[Doc – Subaward Attachment 4, p. 6]

“Could you review your response and add more precise citations?”

Redacted and condensed from the authentic INCLUDE exchange · Feb. 10, 2026

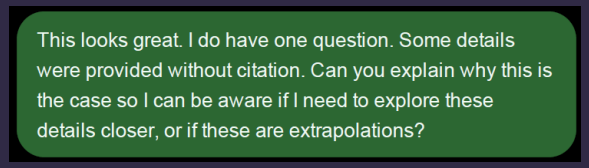
The outputs began to change.

Three documented stages. Not a controlled model benchmark.

AUG 2025 / ORIGIN

PI Checklist mapped into a repeatable extraction.

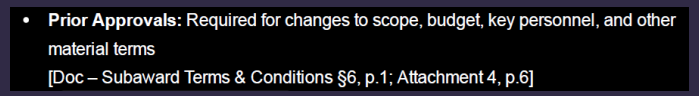
Citations and source tags emerged through human corrections.



FEB 2026 / REFINEMENT

Useful extraction and [Doc] labels.

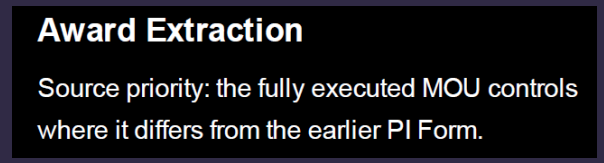
Citation precision still needed repeated requests.



AUG 2026 / MATURITY

Explicit controlling-source priority.

Document conflicts, inferred calculations, and follow-up flags.



Historical examples: BeachSMART → INCLUDE → final project-folder extraction

More information isn’t always better.

The complete extraction and the PI-facing checklist serve different purposes.

AI EXTRACTION

Comprehensive award provisions

Supporting citations and interpretations

Missing information and follow-up issues

HUMAN-CURATED OUTPUT

What the PI needs to know

What finance needs to act on

What stays in the grant file

VERIFICATION → INTERPRETATION → AUDIENCE CALIBRATION

We weren’t just improving a prompt.

We were gradually defining professional behavior.

SCOPE

What belongs in the review?

SOURCES

Which document controls?

EVIDENCE

How can we verify it?

OUTPUT

Who needs to know?

Instructions → Repeatable workflow → Agent architecture

My management experience was more useful

...than my computer science degree.

MANAGEMENT

Expectations
Delegation
Feedback

AI WORKFLOW

Scope
Instructions
Review

A useful analogy—not a claim that AI learns or understands like a human employee.

Don’t just assign a task.

Establish the information environment.

01 INFORMATION SEEKING BEHAVIOR

What is in—and out—of the review?

02 SOURCES PRIORITY

Which documents control when they disagree?

03 SUCCESS

What must the output look like?

04 UNCERTAINTY

When should the agent ask or flag?

An unexpected result is information.

Find the cause before rewriting the instruction.

Unsupported inference

The model supplied a plausible answer without adequate evidence.

Conflicting sources

Multiple documents suggested different answers or priorities.

Instruction conflict

The output followed one rule while neglecting another.

Observe → Diagnose → Refine → Retest

Iteration isn’t just for programmers.

Sometimes the best update is a better SOP.

BEFORE

Reporting requirements

One combined list of obligations.

AFTER

PI / Grants Accountant

Responsibilities organized for the people who act on them.

Same source documents. Better communication and review.

The citation was useful. It still wasn’t enough.

FIRST CORRECTION

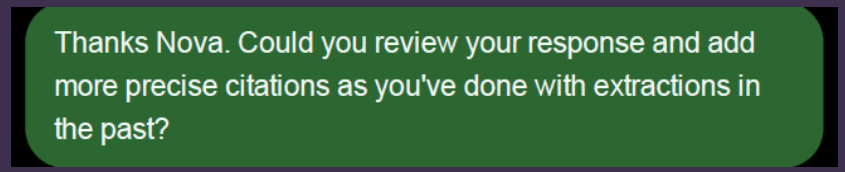
Precise citations added to the extracted grant provisions.

[Doc – Subaward Attachment 4, p. 6]

→

REMAINING GAP

The PI / Grants Accountant reporting summary still lacked precise citations.



Human follow-up → a separately cited reporting section

One revision helped. A second review exposed the next requirement.

Historical example: February 10, 2026 · INCLUDE extraction · condensed and de-identified

Did AI save me work?

Or did it just give me different work?

Accuracy	Are claims supported by the documents?
Consistency	Does the format stay predictable?
Boundaries	Does the system stay within its role?
Human effort	Did verification and correction become easier?

Cross-agent pressure testing can reveal problems—not replace independent verification.

Agent Seed:
We packaged our professional requirements.

A reusable specification became Penny's foundation.

ROLE
Award setup support

SOURCE ORDER
Controlling documents

EVIDENCE
Tags and citations

BOUNDARIES
No rewriting terms

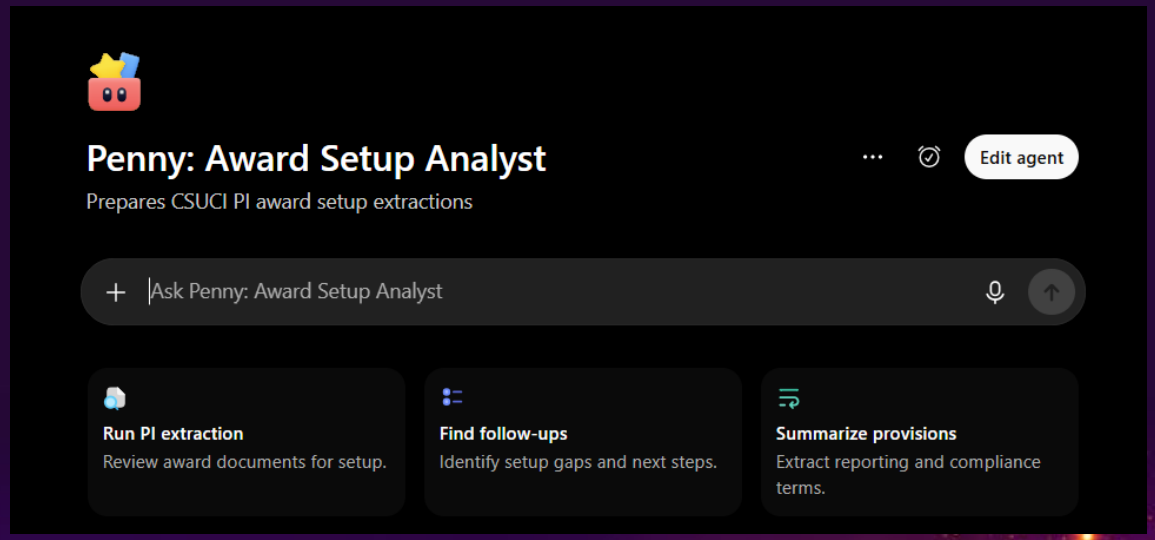
OUTPUTS
PI-ready structure

OVERSIGHT
Human decisions

Meet Penny.

A specialized AI workflow for award setup.

Built from professional requirements



The model is not the architecture.

General capabilities become a specialized workflow through explicit professional expectations.

Language model

General language and reasoning capabilities. A new model can improve performance without changing our professional requirements.

+

Professional architecture

Purpose, source priorities, boundaries, citations, output formats, and human review.

PENNY = CAPABILITIES + A DOCUMENTED WORKFLOW

Professional knowledge, made explicit.

The seed documents how the work should be performed—not merely what answer to generate.

01 **Purpose**

PI-ready award setup

02 **Sources**

Controlling document hierarchy

03 **Evidence**

Source tags + page citations

04 **Boundaries**

No unauthorized decisions

05 **Outputs**

Checklist + detailed analysis

06 **Oversight**

Follow-up and human review

The public agent seed is available in the GitHub toolkit.

The difference is in the evidence trail.

August 27, 2026 • One MOU and one internal form • Two documented workflows

PROJECT-FOLDER EXTRACTION

“The executed MOU controls where it differs from the earlier PI Form.”

Source priority stated; date conflict identified; funding interpreted.

PENNY’S FIRST EXTRACTION

“Files reviewed” • “Controlling award document” • “Execution status”

Also flagged a referenced publisher contract that was not supplied.

More explicit checkpoints—not proof of a model-only improvement.

A claim is easier to review when its source is visible.

The mature workflow already cited sources; Penny made the evidence format more systematic.

PROJECT-FOLDER WORKFLOW

Project end date

Executed MOU controls over the earlier internal form.

MOU § IV.2 / § IV.4, p. 3; PI Form, p. 1

PENNY'S FIRST EXTRACTION

Execution status

Signatures and execution date identified from the governing agreement.

[Doc] (MOU § IV.4, PDF p. 3)

Different findings, not a matched-field accuracy test.

The PI Checklist is the finished communication.

Reconstructed from the completed PSA checklist • identifying fields removed

PRINCIPAL INVESTIGATOR FORM

For acceptance and setup of grants and contracts

Investigator / Program / Award identifiers: [REDACTED]

GRANT PROVISIONS

- Modifications and approvals
- Early termination notice
- Transition responsibilities
- Intellectual property and permissions

REPORTING REQUIREMENTS

PI: annual report and interim communications
Grant accountant: invoicing window, payment method and timing

What the PI needs

A few clear responsibilities.

What stays in the file

The comprehensive cited extraction.

Human verification + selection

One extraction. Multiple professional uses.

The AI report is an analytical resource—not the final communication.

1 **Source packet**
Agreement + internal forms

→

2 **Penny's report**
Detailed, cited extraction

→

3 **Human review**
Verify • select • contextualize

→

4 **PI Checklist**
Concise, meeting-ready

PI-facing document

Full extraction retained in grant file

Up to 3 hours → about 30 minutes

Personal experience with some complex awards; includes human review. Not a controlled benchmark.

Penny can flag the problem.

The institution must evaluate and resolve it.

! Conflicting PI information

Notice and internal form identify different people

! Conflicting agency identification

Header differs from governing terms

! Inconsistent obligation fields

Different award boxes show different amounts

! Missing incorporated application

Scope and budget documents not supplied

AI observation ≠ institutional determination

Penny was the proof of concept.

The same design principles can support other research administration tasks.

Scope • Sources • Evidence • Boundaries • Outputs • Oversight

“I’m actually a pre-award analyst.”

Different agents. Different responsibilities.

A growing collection of specialized experiments. Not an autonomous department.

SCOUT

Funding opportunities and proposal readiness

QUILL

Proposal data entry assistant

PENNY

Award setup and agreement interpretation

MAGPIE

Legacy grant database knowledge

The work changes. The design principles remain.

A repeatable architecture can be adapted to different professional tasks.

SCOPE	SOURCE PRIORITIES	BOUNDARIES
EVIDENCE	OUTPUT FORMAT	HUMAN OVERSIGHT

PENNY / Award agreements

SCOUT / Funding solicitations

From AI disclosure to AI attribution

Not merely “I used AI”—explain what it contributed and who made the decision.

Disclosure

“I used AI to review this budget.”

Attribution

“I verified the budget; Nova identified a possible narrative inconsistency that I evaluated and flagged.”

Attribution is not verification. Cite the underlying evidence. Retain human accountability.

Don’t just share the result.

Share how it was built.

RESEARCH ADMINISTRATION AI TOOLKIT

Penny + Scout agent seeds

Historical guidelines • Transparent AI Guide

github.com/Fabledskye/research-admin-ai-toolkit

Start with a problem you already understand.

A practical, low-risk starting point for the people in this room.

01 Choose one recurring cognitive task.

02 Identify authoritative, permitted sources.

03 Describe what success looks like.

04 Set boundaries and rules for uncertainty.

05 Test, verify, refine—and measure human effort.

The breakthrough wasn’t teaching AI to know everything.

It was learning how to work with what it knows—and what it doesn’t.

Begin with something you already know.

The conversation doesn’t end here.

Connect with Fable or explore the public toolkit.

CONNECT WITH ME

Fable Mangold

Grants Analyst | CSU Channel Islands

LinkedIn / linkedin.com/in/fable-mangold/

fable.mangold@csuci.edu

EXPLORE THE TOOLKIT

Penny & Scout

Agent seeds + guides

github.com/Fabledskye/research-admin-ai-toolkit

Begin with something you already know.